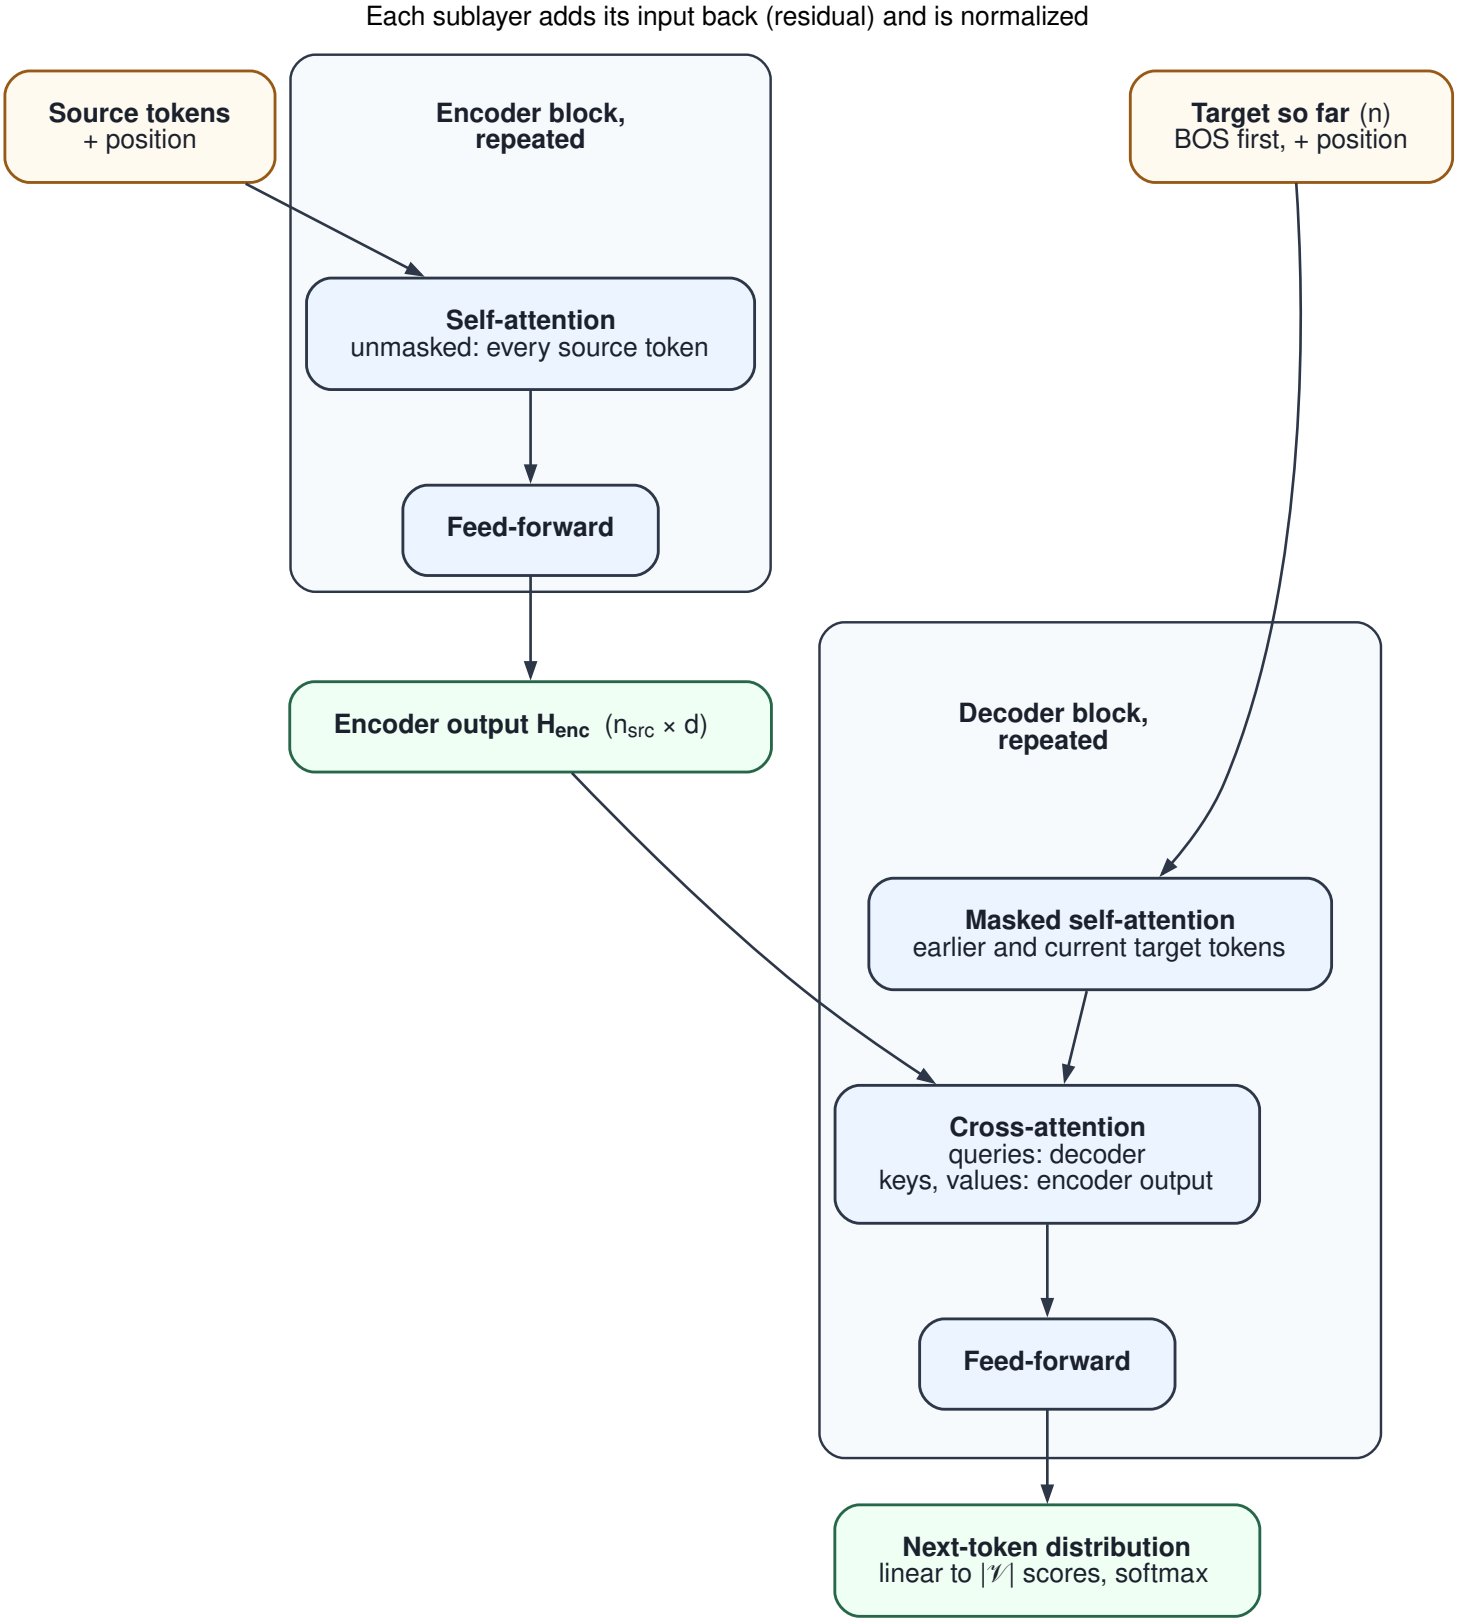
Each sublayer adds its input back (residual) and is normalized
Source tokens
+ position
Encoder block,
repeated
Self-attention
unmasked: every source token
Feed-forward
Encoder output $H_{enc}$ ($n_{src}$ × d)
Target so far (n)
BOS first, + position
Decoder block,
repeated
Masked self-attention
earlier and current target tokens
Cross-attention
queries: decoder
keys, values: encoder output
Feed-forward
Next-token distribution
linear to |𝒱| scores, softmax