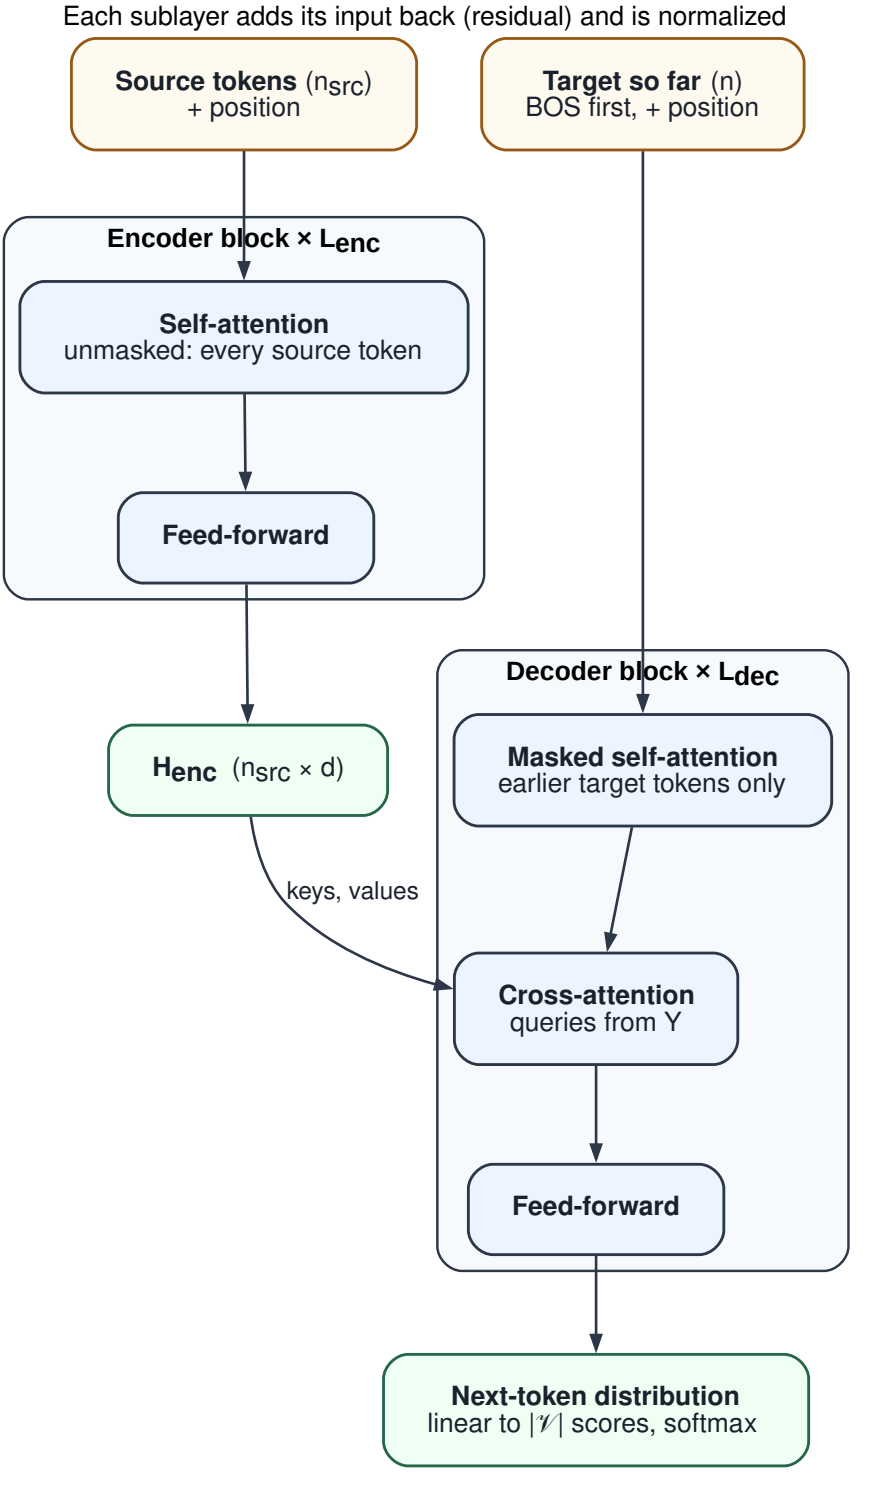
Each sublayer adds its input back (residual) and is normalized
Source tokens (nsrc) + position
Target so far (n) BOS first, + position
Encoder block × Lenc
Self-attention unmasked: every source token
Feed-forward
Henc (nsrc × d)
Decoder block × Ldec
Masked self-attention earlier target tokens only
keys, values
Cross-attention queries from Y
Feed-forward
Next-token distribution linear to |𝒱| scores, softmax