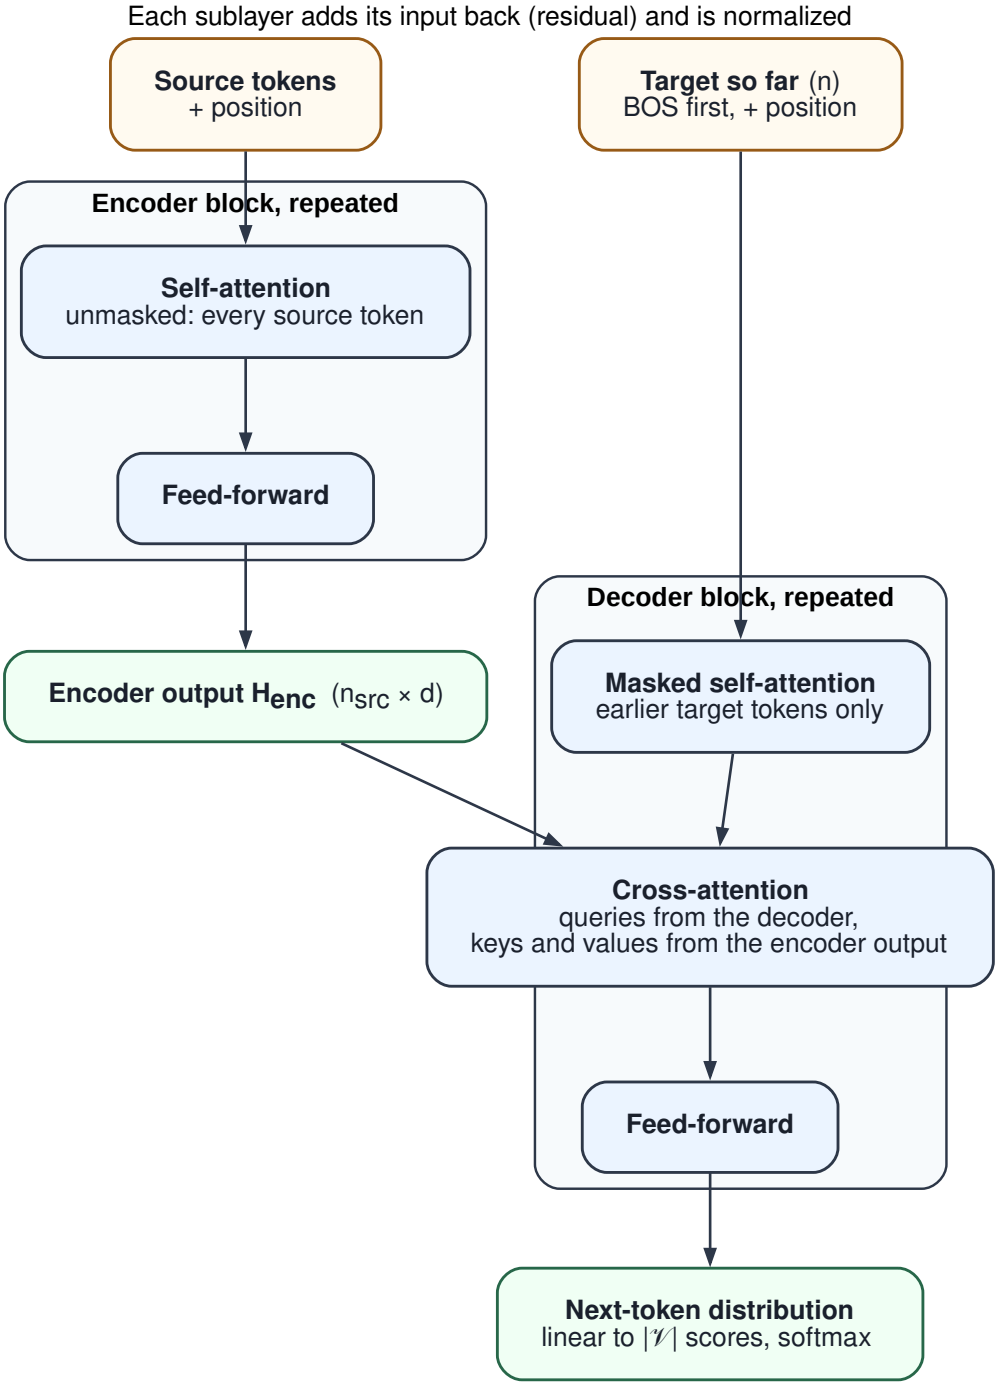
Each sublayer adds its input back (residual) and is normalized
Source tokens
+ position
Target so far (n)
BOS first, + position
Encoder block, repeated
Self-attention
unmasked: every source token
Feed-forward
Encoder output H_enc (n_src × d)
Decoder block, repeated
Masked self-attention
earlier target tokens only
Cross-attention
queries from the decoder,
keys and values from the encoder output
Feed-forward
Next-token distribution
linear to |𝒱| scores, softmax