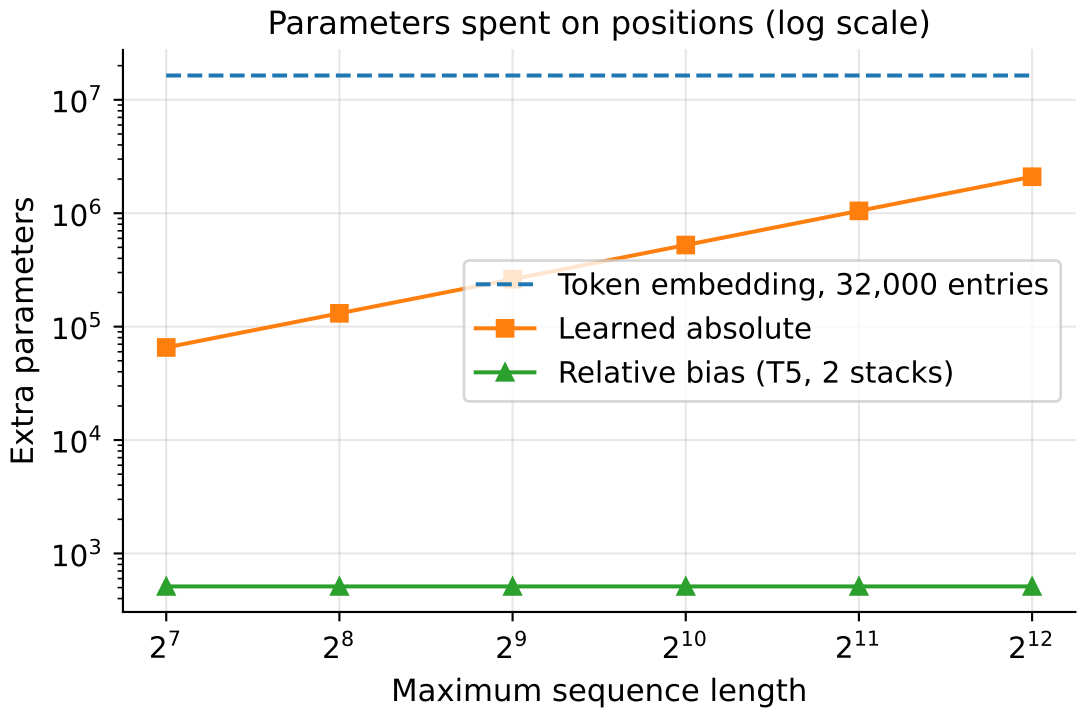
Parameters spent on positions (log scale)
Extra parameters
10^7
10^6
10^5
10^4
10^3
Token embedding, 32,000 entries
Learned absolute
Relative bias (T5, 2 stacks)
2^7
2^8
2^9
2^{10}
2^{11}
2^{12}
Maximum sequence length