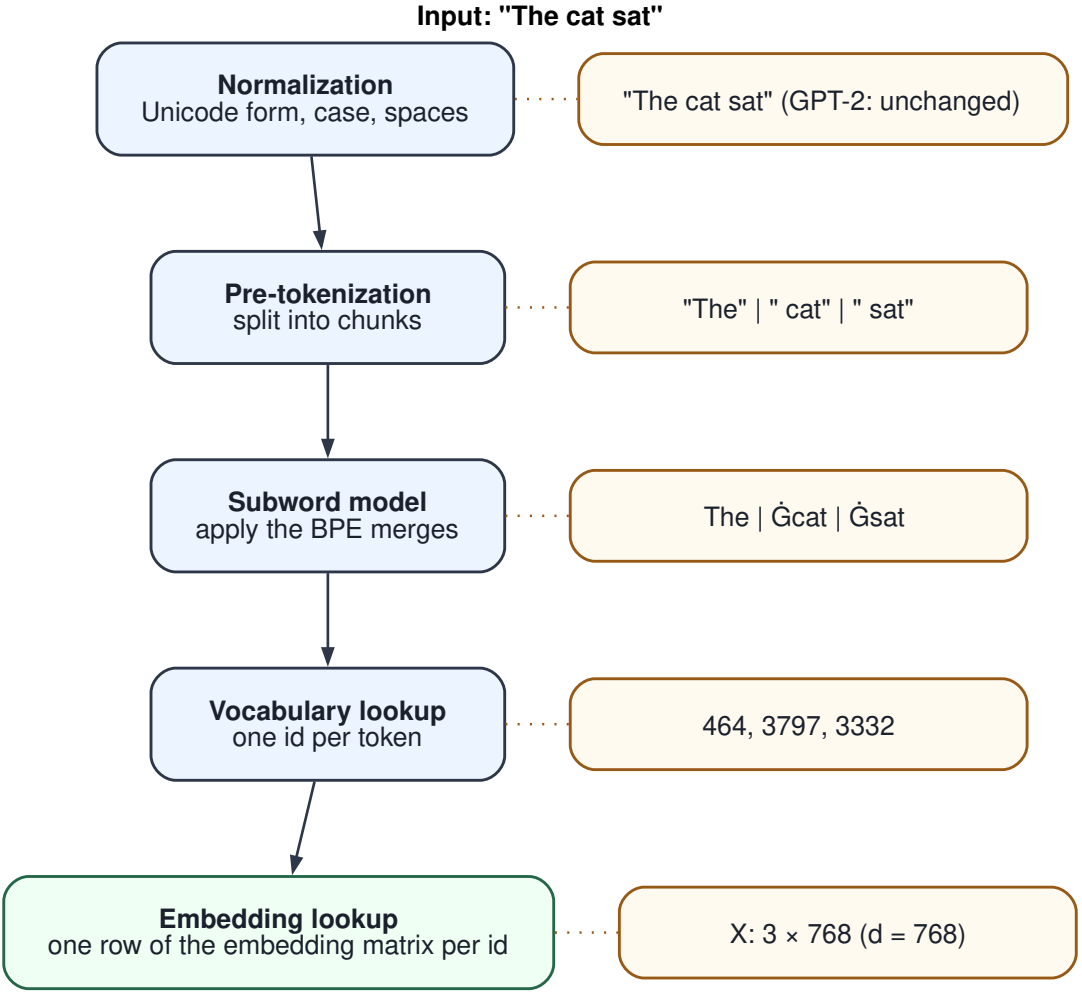
Input: "The cat sat"
Normalization
Unicode form, case, spaces
"The cat sat" (GPT-2: unchanged)
Pre-tokenization
split into chunks
"The" | " cat" | " sat"
Subword model
apply the BPE merges
The | Ġcat | Ġsat
Vocabulary lookup
one id per token
464, 3797, 3332
Embedding lookup
one row of the embedding matrix per id
X: 3 × 768 (d = 768)