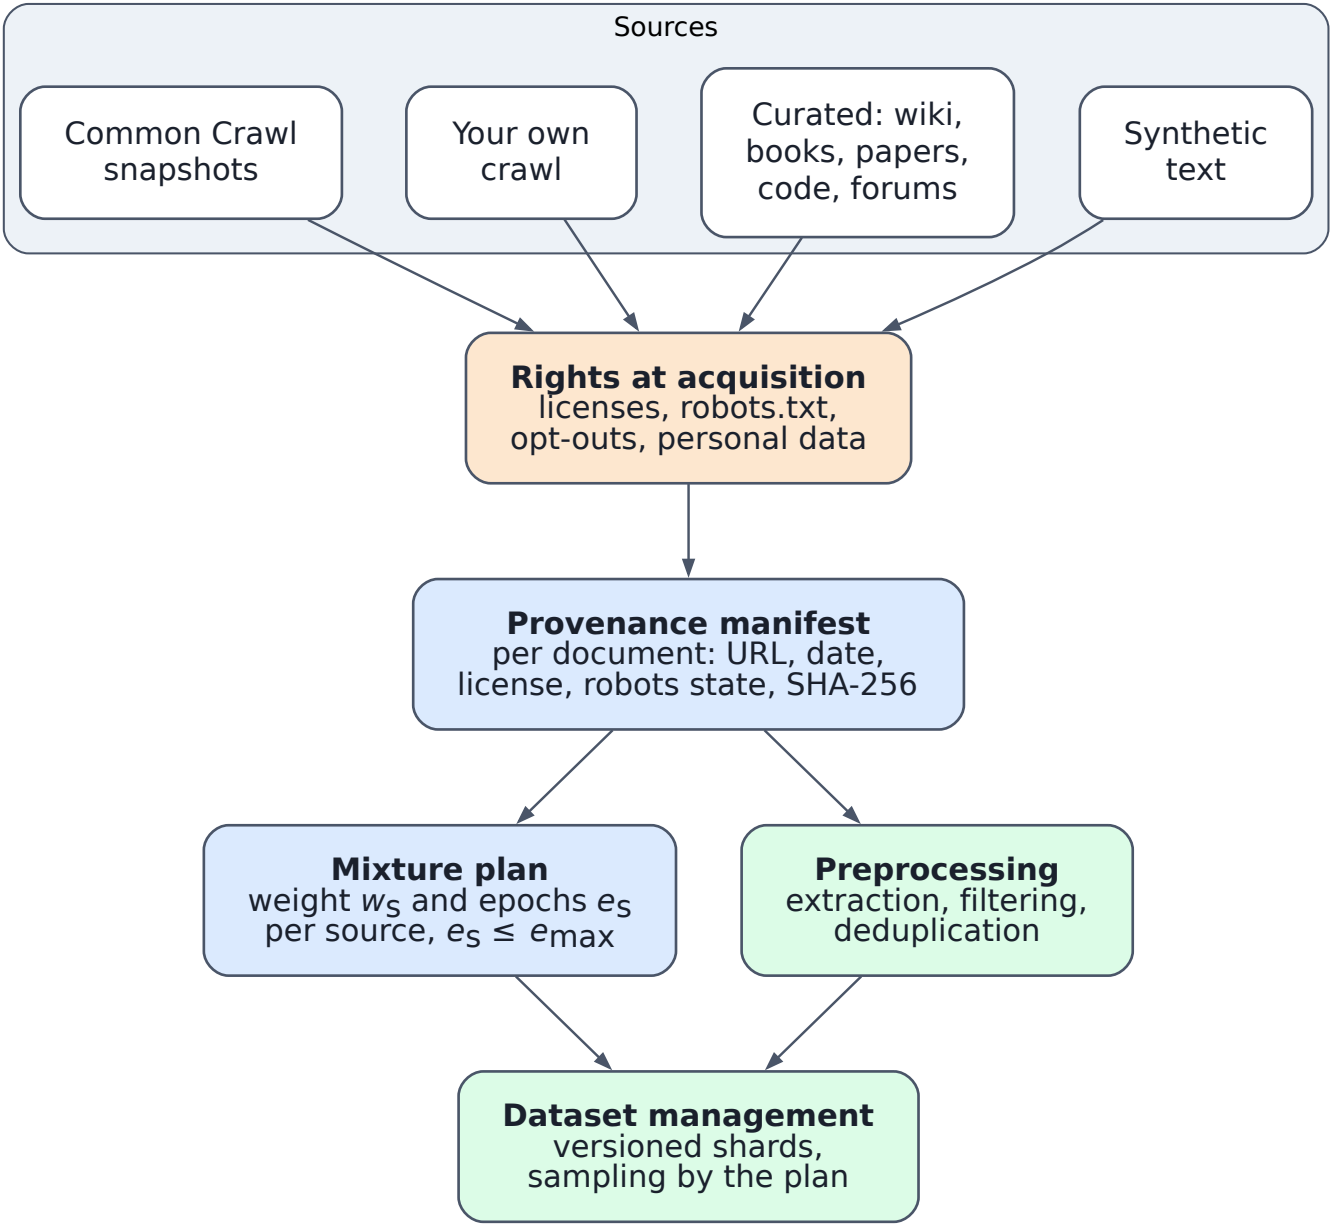
Sources
Common Crawl snapshots
Your own crawl
Curated: wiki, books, papers, code, forums
Synthetic text
Rights at acquisition
licenses, robots.txt, opt-outs, personal data
Provenance manifest
per document: URL, date, license, robots state, SHA-256
Mixture plan
weight w_S and epochs e_S per source, $e_S \leq e_{max}$
Preprocessing
extraction, filtering, deduplication
Dataset management
versioned shards, sampling by the plan