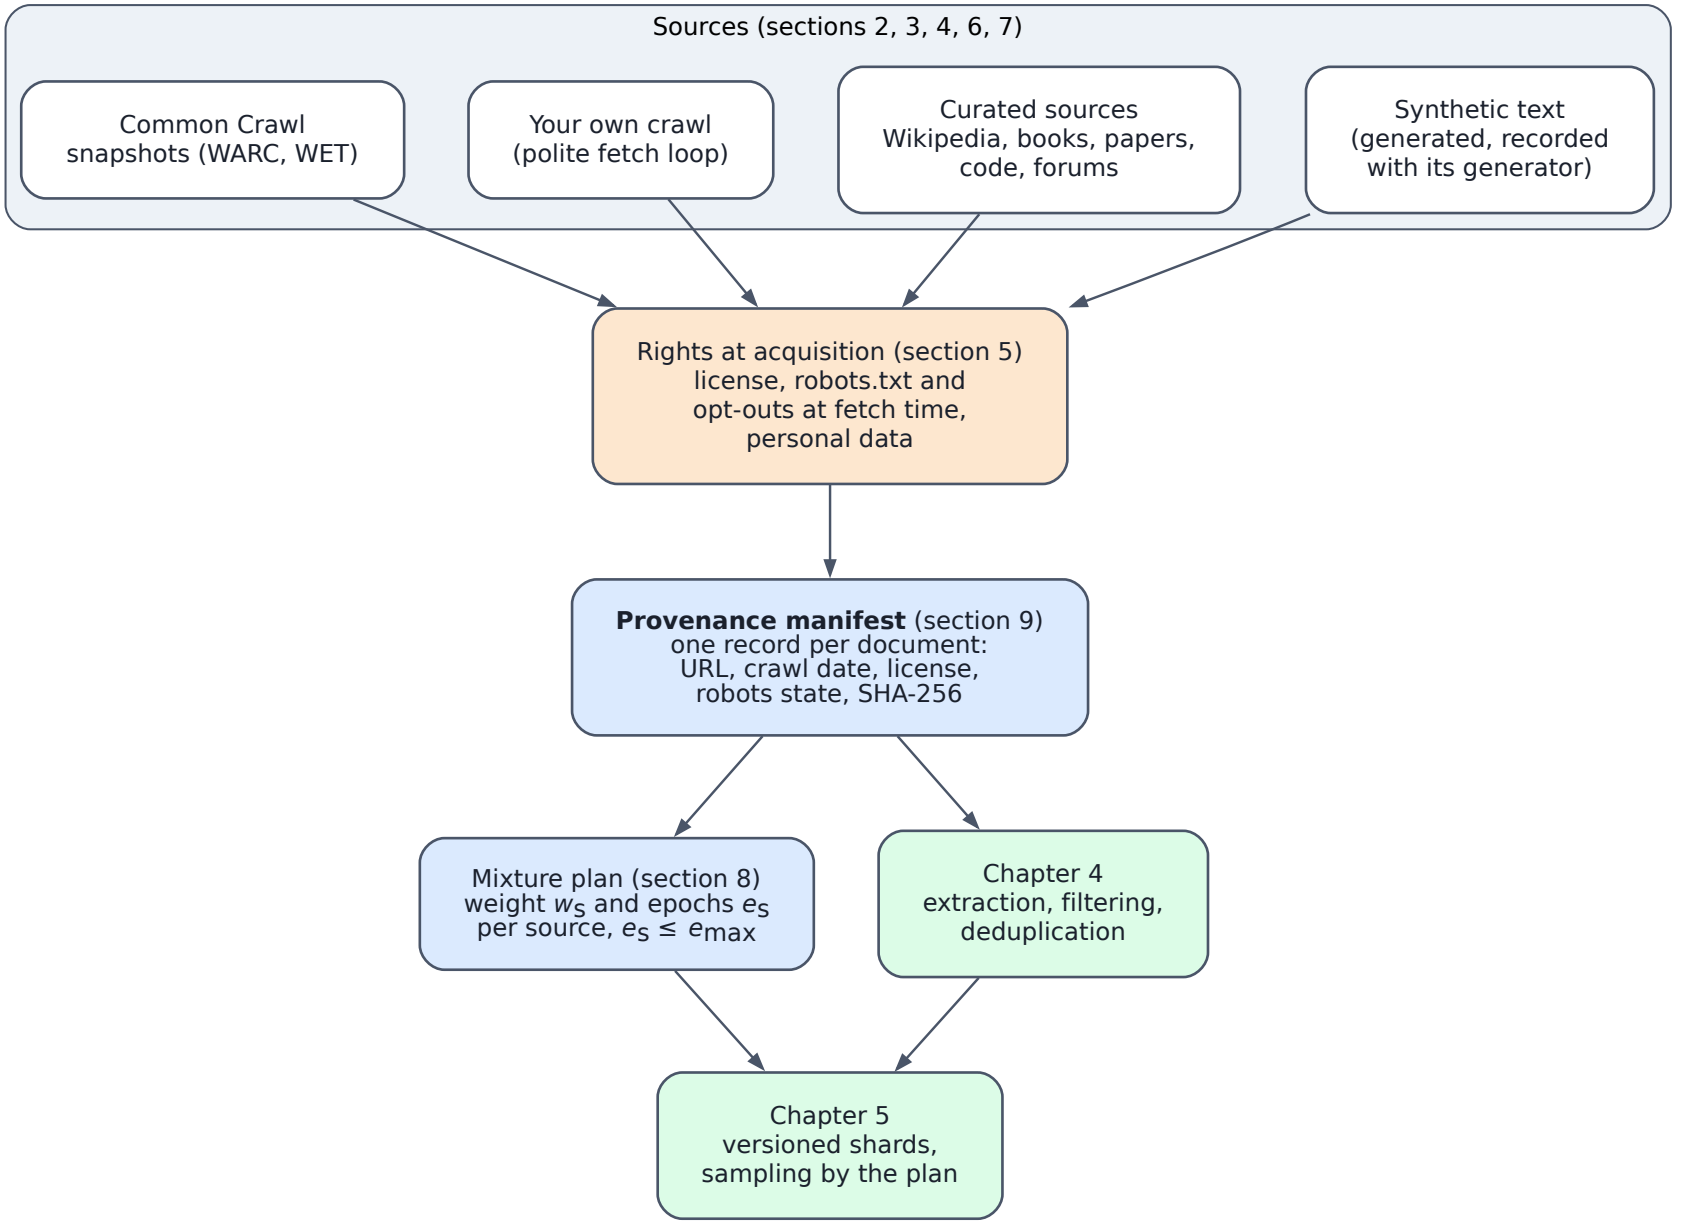
Sources (sections 2, 3, 4, 6, 7)
Common Crawl snapshots (WARC, WET)
Your own crawl (polite fetch loop)
Curated sources Wikipedia, books, papers, code, forums
Synthetic text (generated, recorded with its generator)
Rights at acquisition (section 5) license, robots.txt and opt-outs at fetch time, personal data
Provenance manifest (section 9) one record per document: URL, crawl date, license, robots state, SHA-256
Mixture plan (section 8) weight w_S and epochs e_S per source, $e_S \le e_{max}$
Chapter 4 extraction, filtering, deduplication
Chapter 5 versioned shards, sampling by the plan