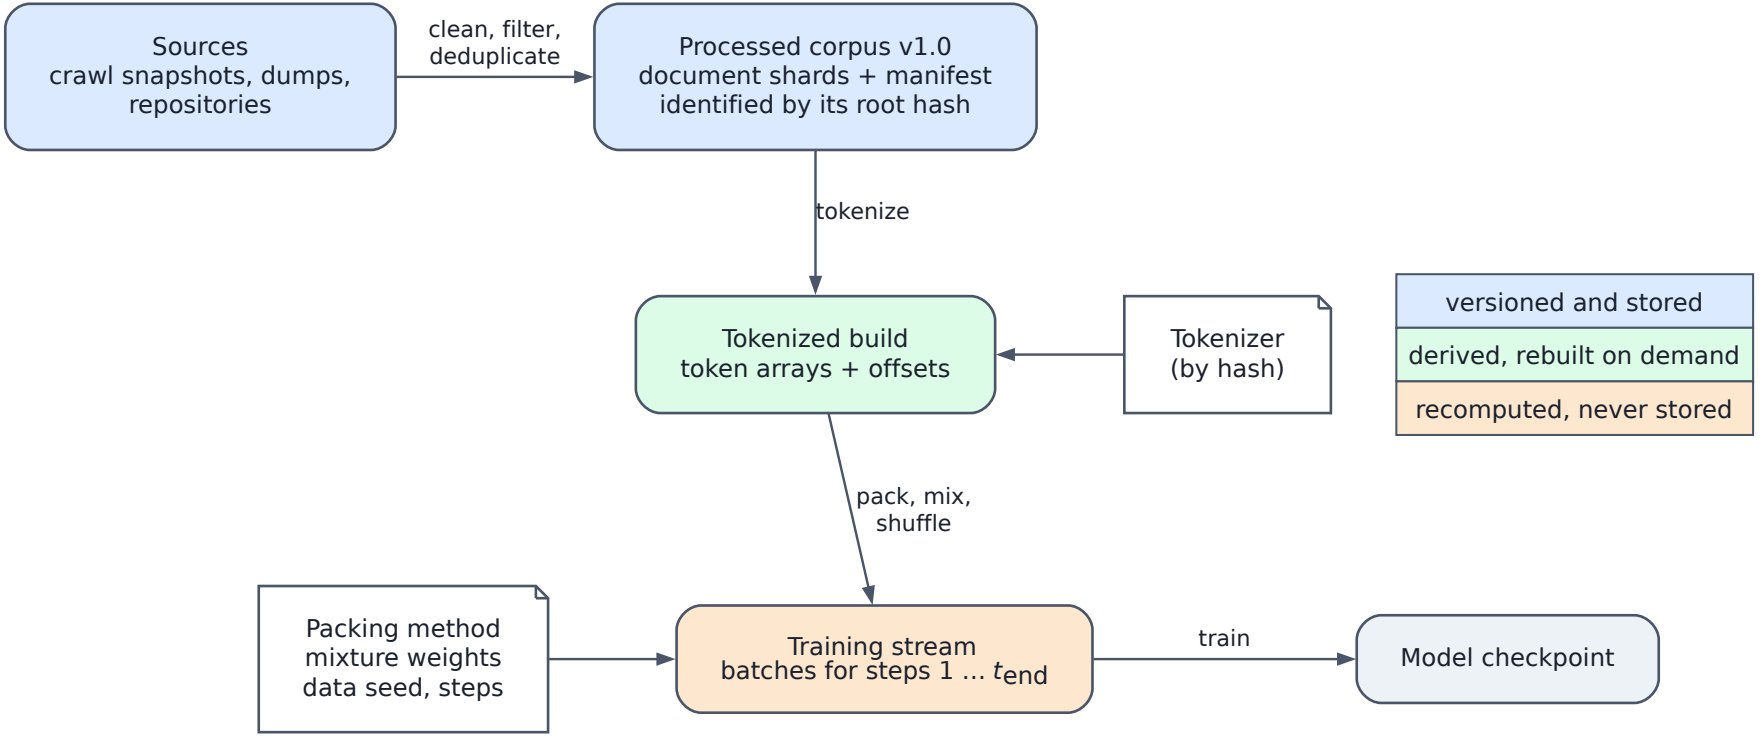
Sources
crawl snapshots, dumps, repositories
clean, filter, deduplicate
Processed corpus v1.0
document shards + manifest
identified by its root hash
tokenize
Tokenized build
token arrays + offsets
Tokenizer
(by hash)
versioned and stored
derived, rebuilt on demand
recomputed, never stored
pack, mix, shuffle
Packing method
mixture weights
data seed, steps
Training stream
batches for steps 1 ... $t_{\text{end}}$
train
Model checkpoint