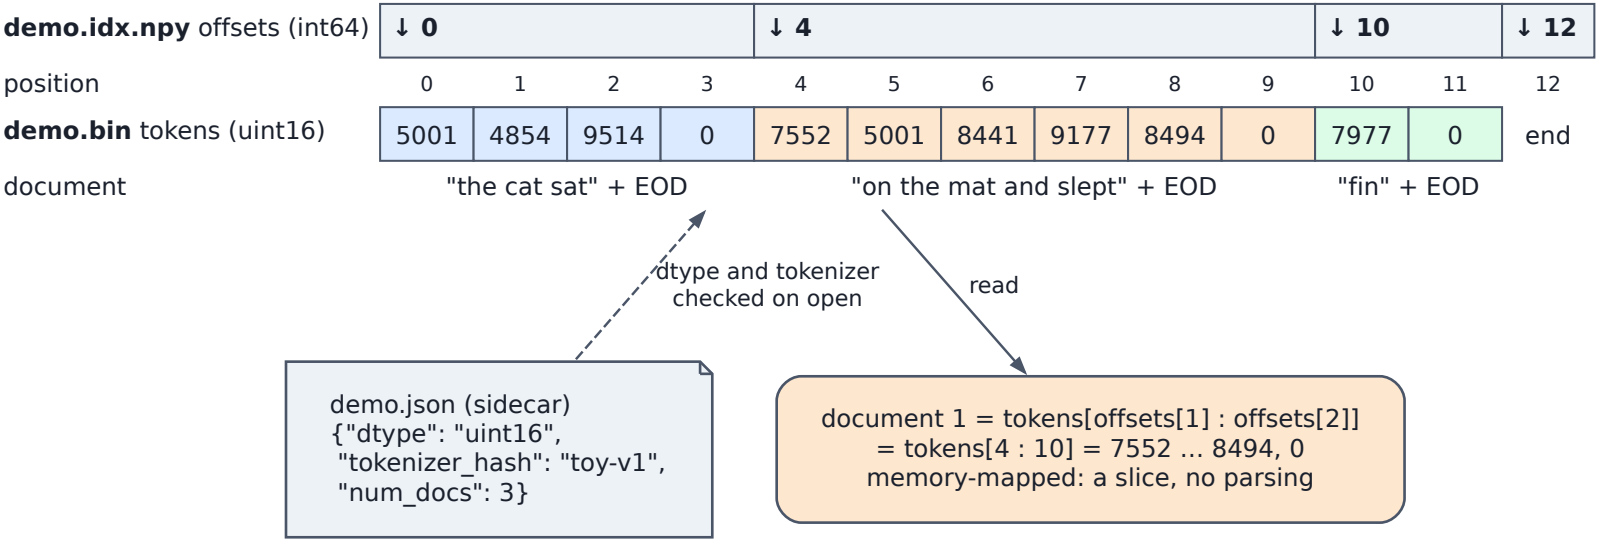
demo.idx.npy offsets (int64)
↓ 0
↓ 4
↓ 10
↓ 12
position
0 1 2 3 4 5 6 7 8 9 10 11 12
demo.bin tokens (uint16)
5001 4854 9514 0 7552 5001 8441 9177 8494 0 7977 0
end
document
"the cat sat" + EOD
"on the mat and slept" + EOD
"fin" + EOD
dtype and tokenizer checked on open
read
demo.json (sidecar)
{"dtype": "uint16",
"tokenizer_hash": "toy-v1",
"num_docs": 3}
document 1 = tokens[offsets[1] : offsets[2]]
= tokens[4 : 10] = 7552 ... 8494, 0
memory-mapped: a slice, no parsing