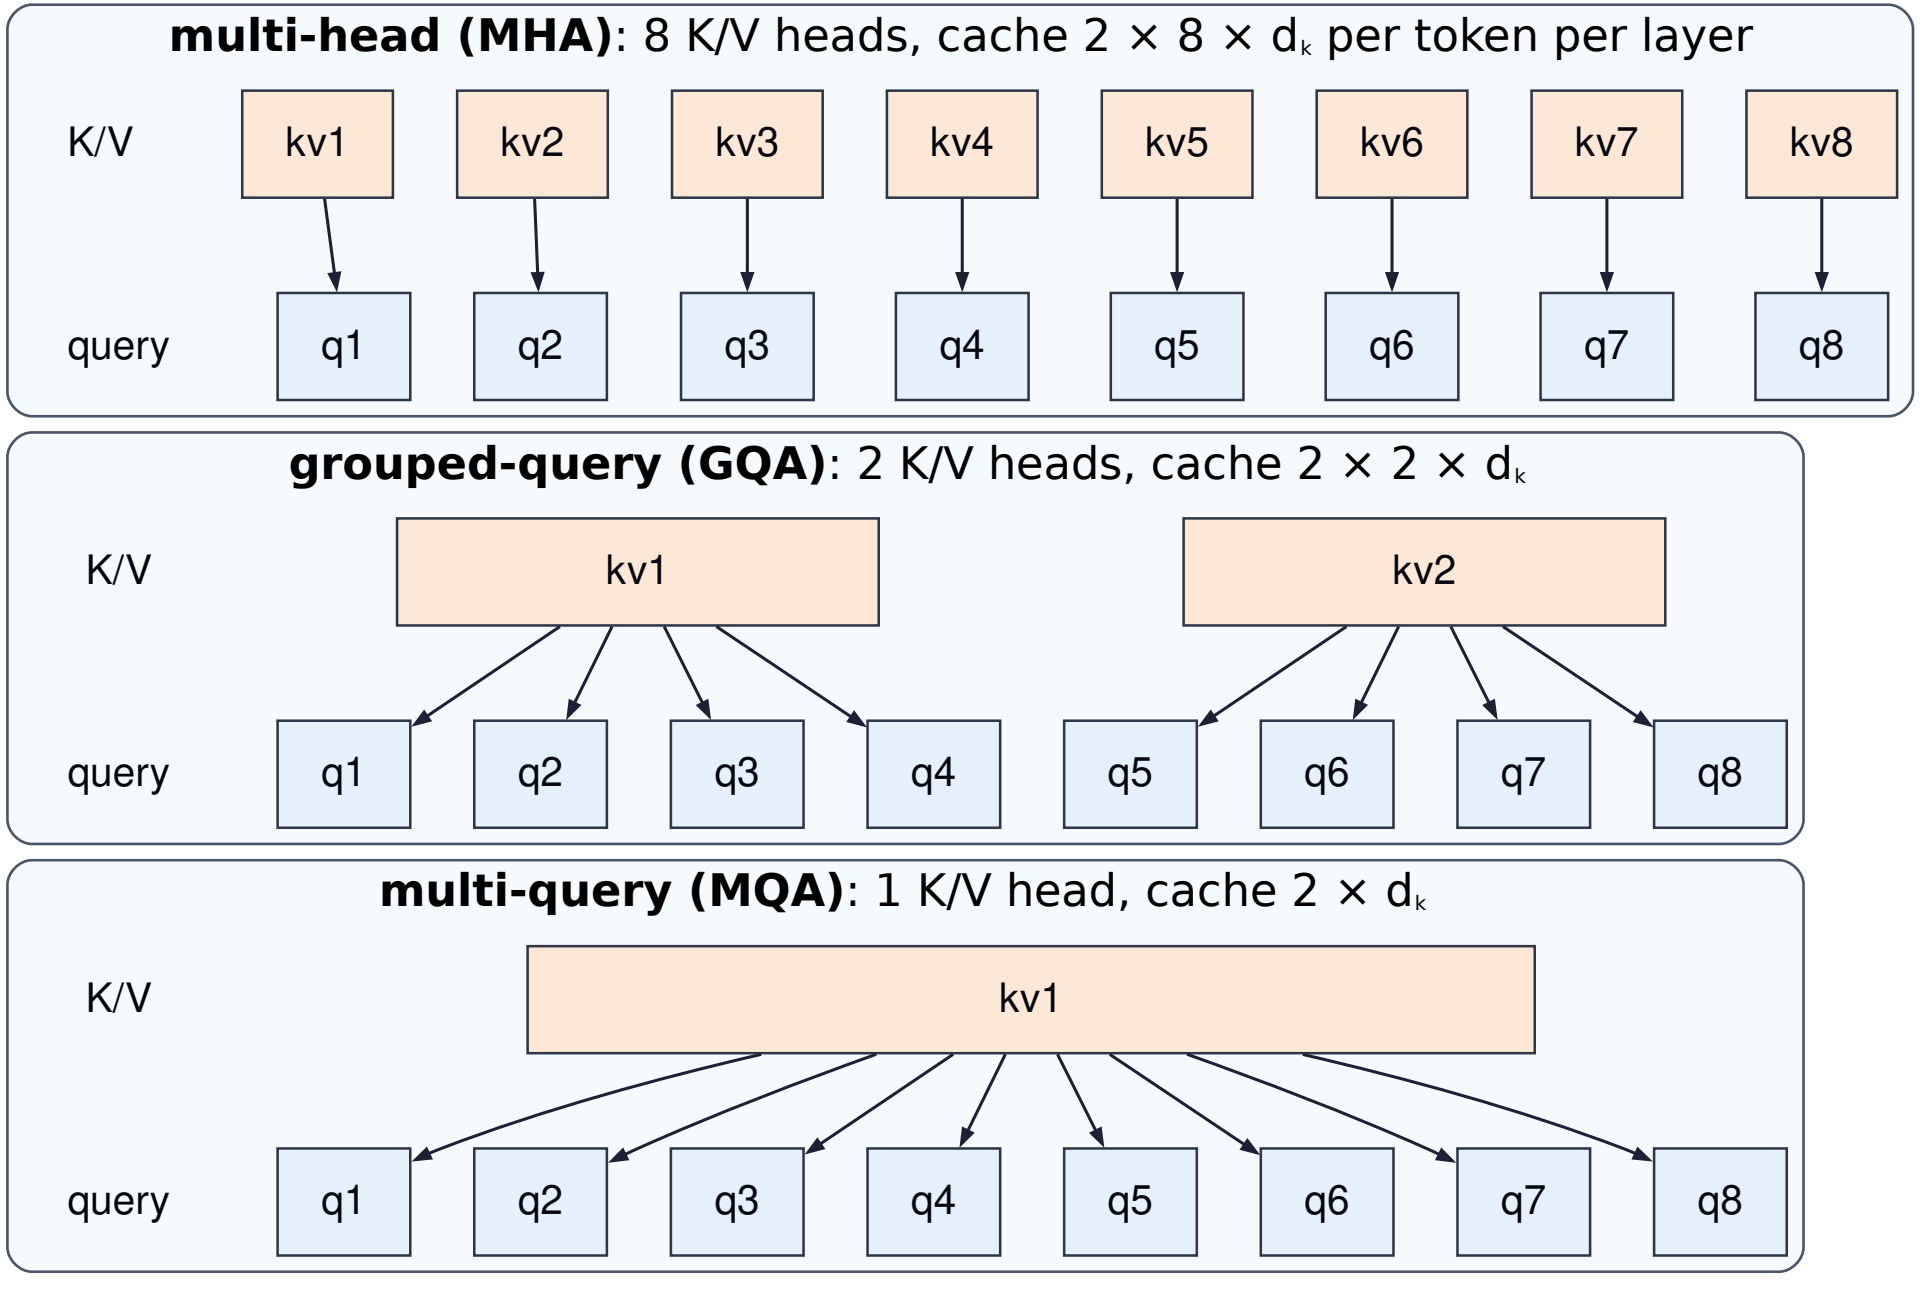
multi-head (MHA): 8 K/V heads, cache 2 × 8 × d_k per token per layer
K/V
kv1
kv2
kv3
kv4
kv5
kv6
kv7
kv8
query
q1
q2
q3
q4
q5
q6
q7
q8
grouped-query (GQA): 2 K/V heads, cache 2 × 2 × d_k
K/V
kv1
kv2
query
q1
q2
q3
q4
q5
q6
q7
q8
multi-query (MQA): 1 K/V head, cache 2 × d_k
K/V
kv1
query
q1
q2
q3
q4
q5
q6
q7
q8