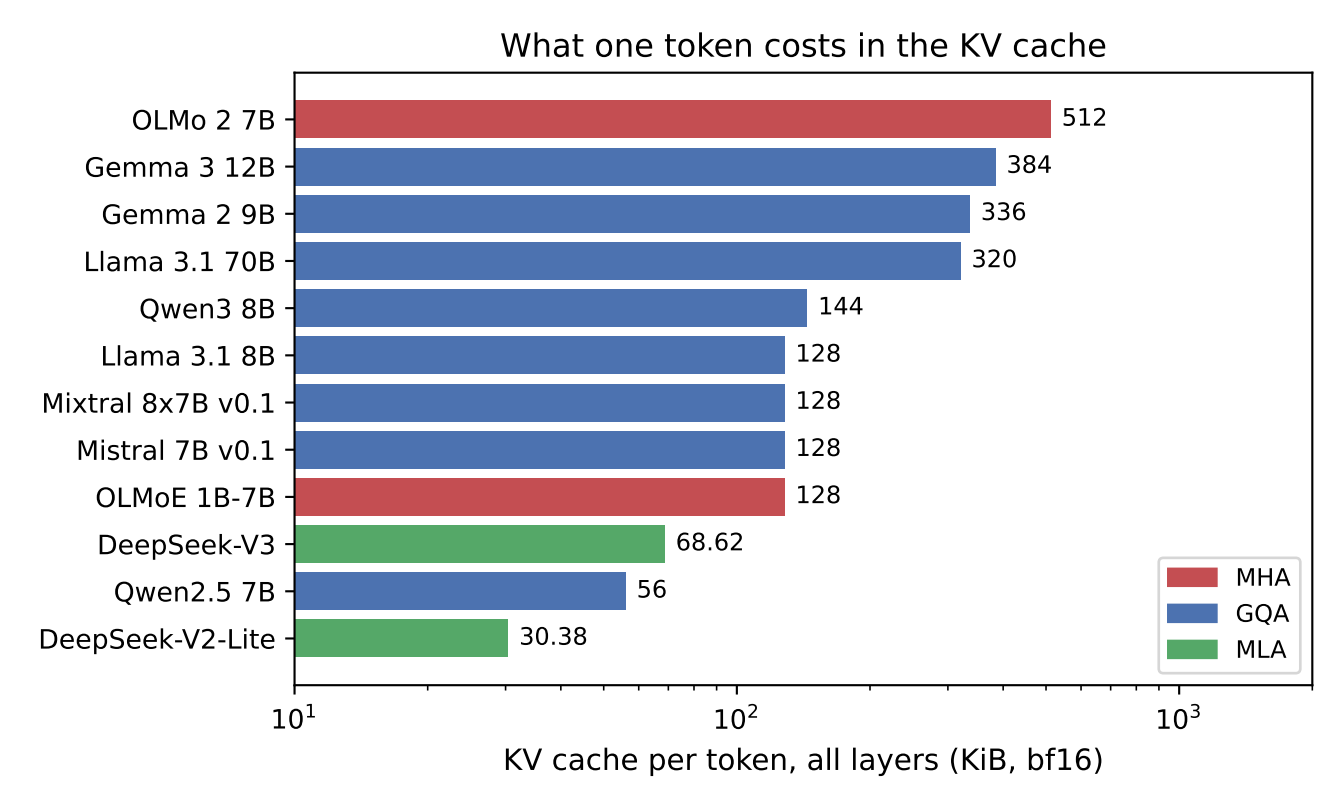
What one token costs in the KV cache
OLMo 2 7B
Gemma 3 12B
Gemma 2 9B
Llama 3.1 70B
Qwen3 8B
Llama 3.1 8B
Mixtral 8x7B v0.1
Mistral 7B v0.1
OLMoE 1B-7B
DeepSeek-V3
Qwen2.5 7B
DeepSeek-V2-Lite
512
384
336
320
144
128
128
128
128
68.62
56
30.38
MHA
GQA
MLA
10^1
10^2
10^3
KV cache per token, all layers (KiB, bf16)