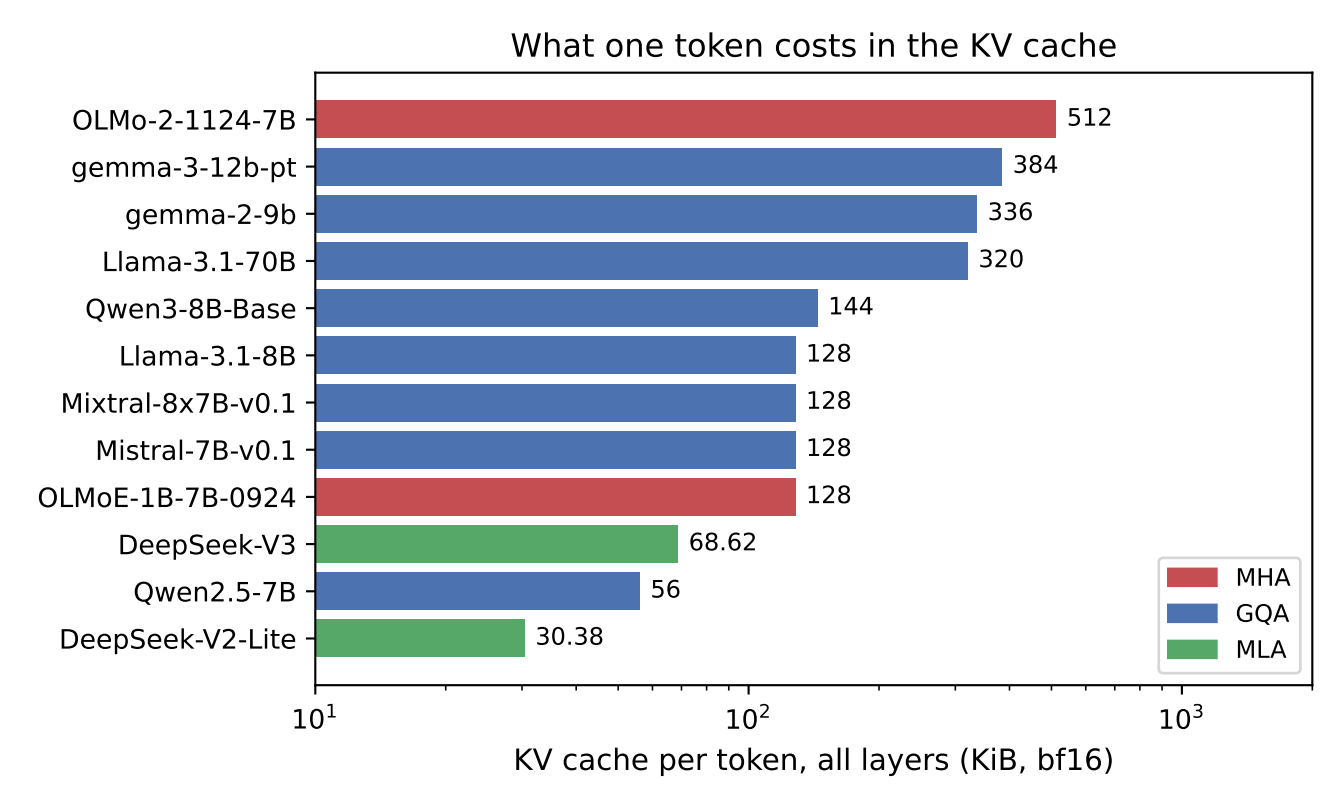
What one token costs in the KV cache
OLMo-2-1124-7B
gemma-3-12b-pt
gemma-2-9b
Llama-3.1-70B
Qwen3-8B-Base
Llama-3.1-8B
Mixtral-8x7B-v0.1
Mistral-7B-v0.1
OLMoE-1B-7B-0924
DeepSeek-V3
Qwen2.5-7B
DeepSeek-V2-Lite
512
384
336
320
144
128
128
128
128
68.62
56
30.38
MHA
GQA
MLA
10^1
10^2
10^3
KV cache per token, all layers (KiB, bf16)