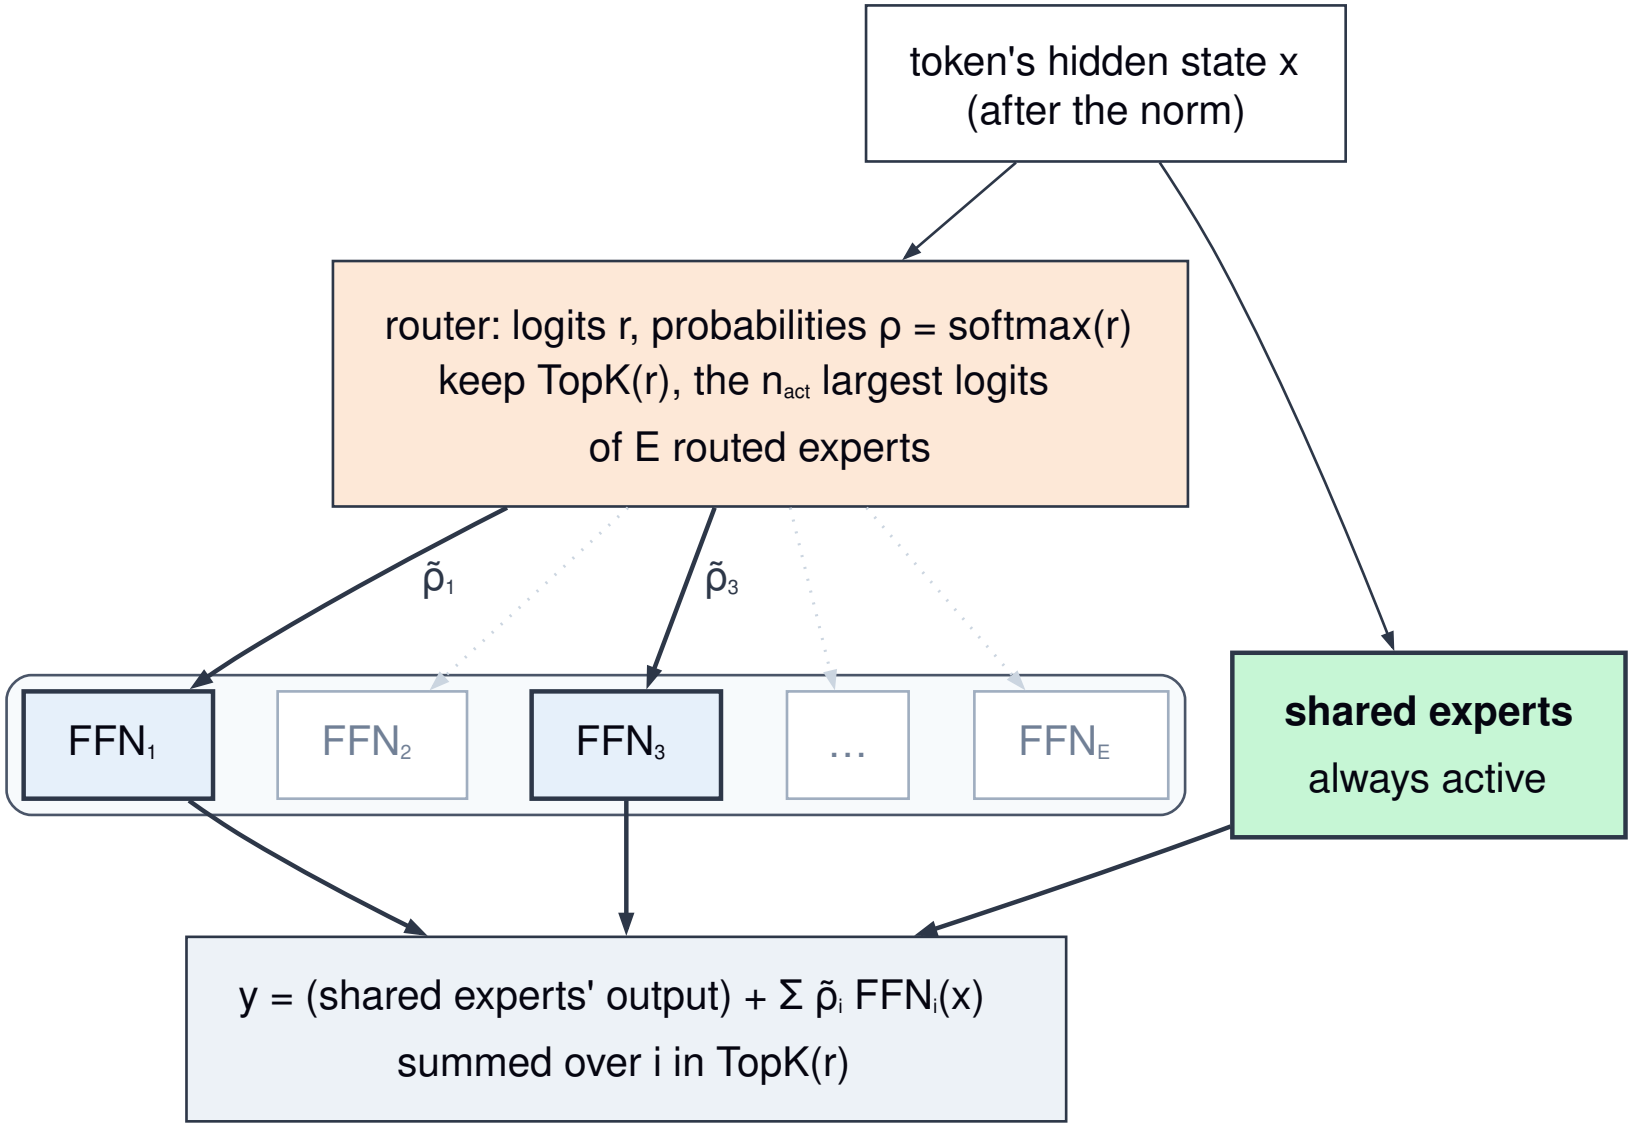
token's hidden state x
(after the norm)
router: logits r, probabilities ρ = softmax(r)
keep TopK(r), the $n_{act}$ largest logits
of E routed experts
$\tilde{\rho}_1$
$\tilde{\rho}_3$
$FFN_1$
$FFN_2$
$FFN_3$
...
$FFN_E$
shared experts
always active
y = (shared experts' output) + Σ $\tilde{\rho}_i$ $FFN_i(x)$
summed over i in TopK(r)