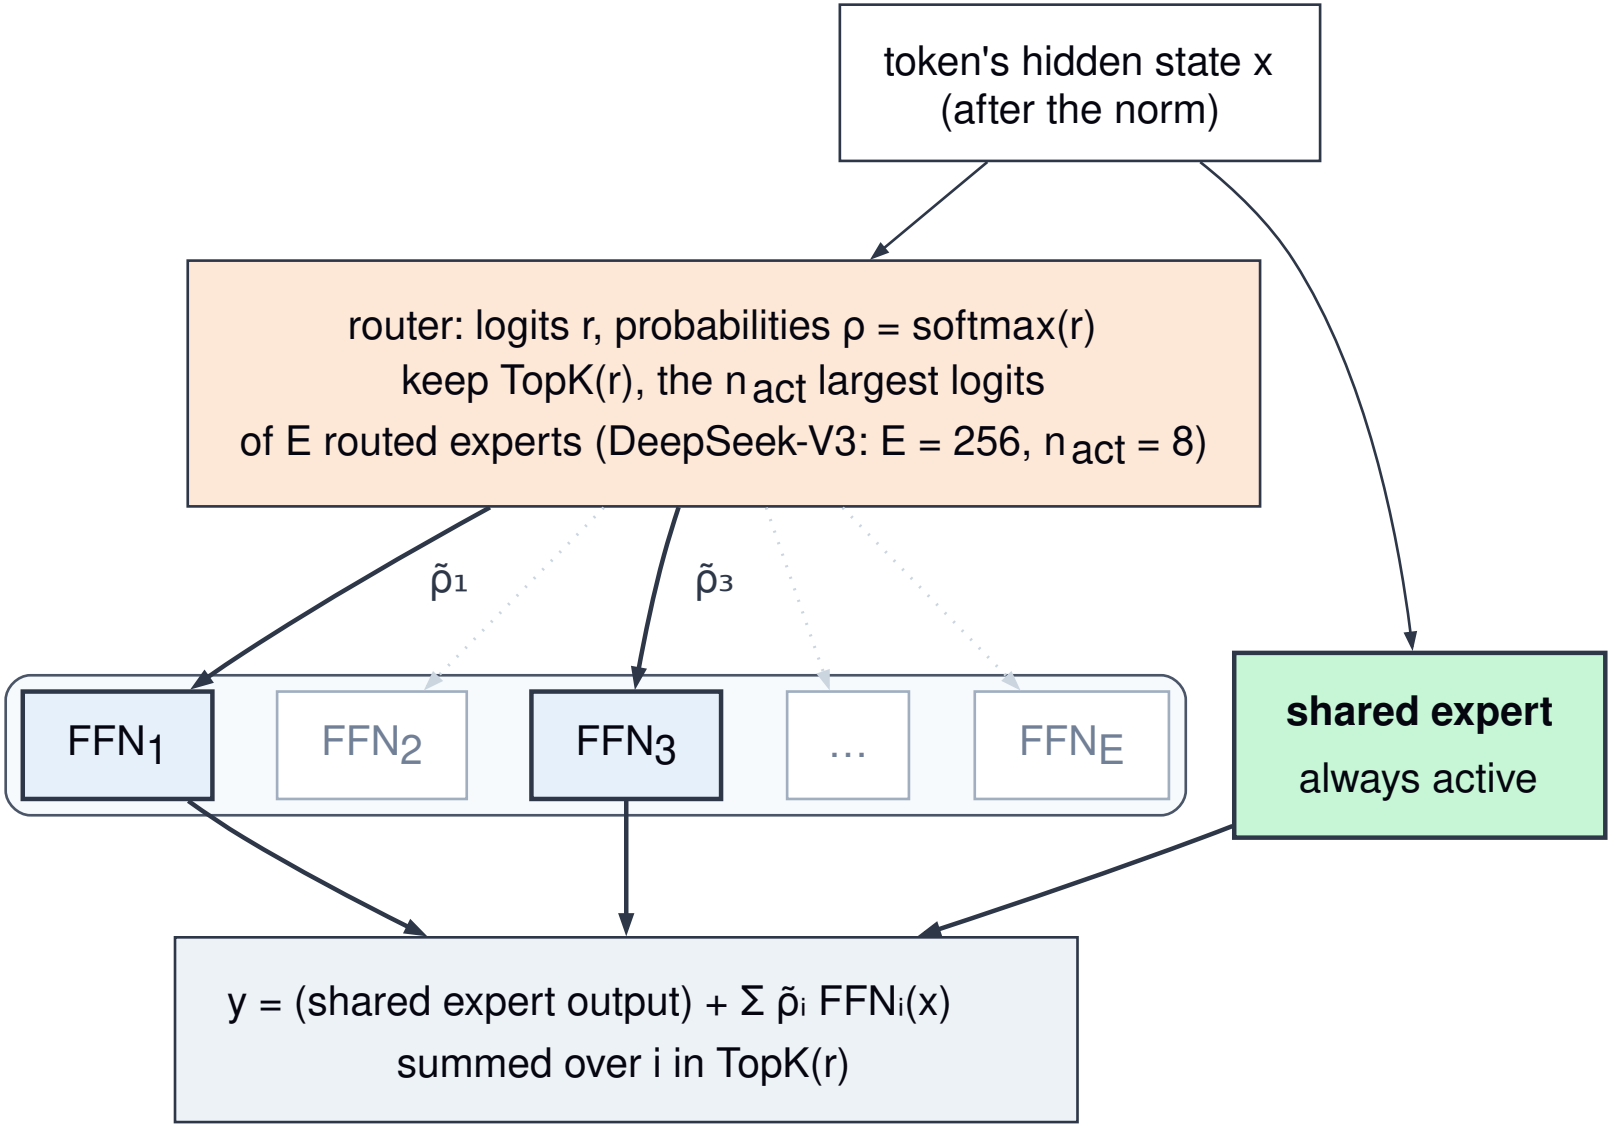
token's hidden state x
(after the norm)
router: logits r, probabilities ρ = softmax(r)
keep TopK(r), the n act largest logits
of E routed experts (DeepSeek-V3: E = 256, n act = 8)
ρ̃1
ρ̃3
FFN1
FFN2
FFN3
...
FFNE
shared expert
always active
y = (shared expert output) + Σ ρ̃i FFNi(x)
summed over i in TopK(r)