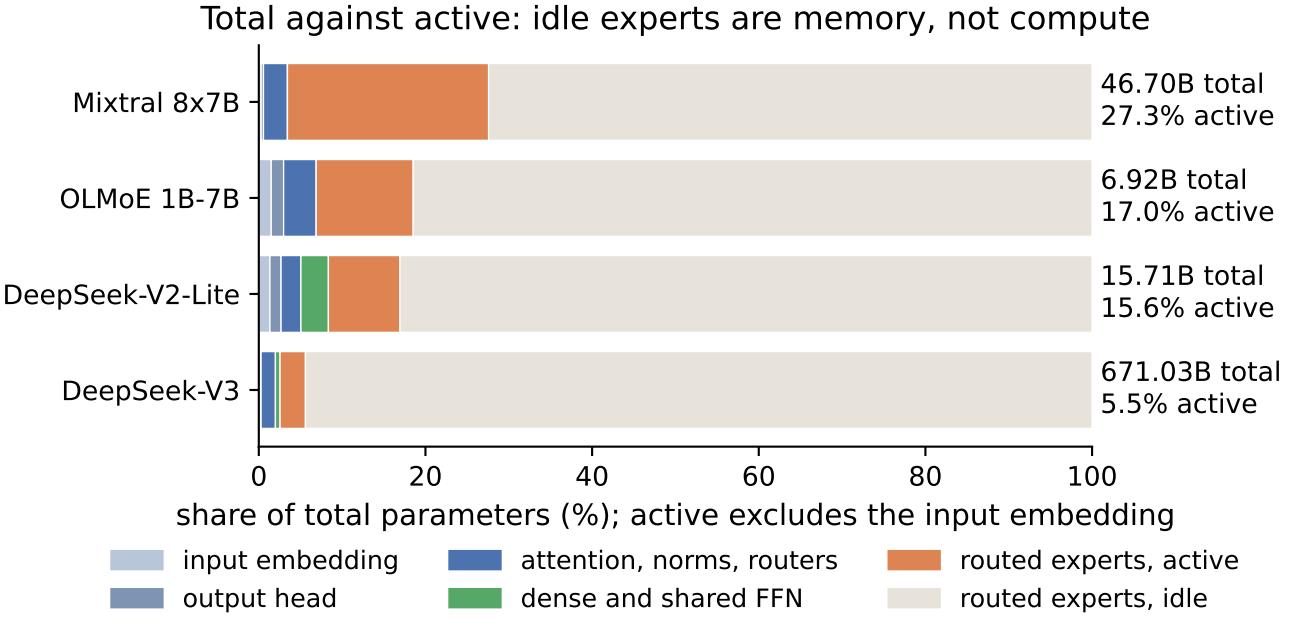
Total against active: idle experts are memory, not compute
Mixtral 8x7B
OLMoE 1B-7B
DeepSeek-V2-Lite
DeepSeek-V3
46.70B total
27.3% active
6.92B total
17.0% active
15.71B total
15.6% active
671.03B total
5.5% active
0
20
40
60
80
100
share of total parameters (%); active excludes the input embedding
input embedding
output head
attention, norms, routers
dense and shared FFN
routed experts, active
routed experts, idle